



Master Individual Project

LNI Reference Checker

Automated Academic Reference Verification and Validation

Submitted by:	Mithila Maruti Prabhu
First examiner:	Prof. Dr. Jürgen Jung
Date of start:	14.04.2026
Date of submission:	15.09.2026

Statement

I confirm that I have written this project report on my own. No other sources were used except those referenced. Content which is taken literally or analogously from published or unpublished sources is identified as such. The drawings or figures of this work have been created by myself or are provided with an appropriate reference. This work has not been submitted in the same or similar form or to any other examination board.

space1cm

Date, Signature of the Student: _____

Contents

1	Introduction	6
1.1	Problem Statement and Context	6
1.1.1	Fabricated and Hallucinated References	6
1.1.2	Metadata Inconsistencies and Transcription Errors	7
1.1.3	Incomplete Bibliography Entries	7
1.1.4	Format Inconsistency and Standardization Challenges	7
1.1.5	Manual Verification Burden and Inconsistency	8
1.2	Motivation and Project Vision	9
1.3	Project Scope and Objectives	10
1.3.1	Primary Objectives	10
1.3.2	In-Scope Elements	10
1.3.3	Out-of-Scope Elements	10
2	Terminology and Concepts	12
2.1	Bibliographic and Citation Terms	12
2.1.1	Citation vs. Bibliography Entry	12
2.1.2	LNI Citation Key Format	12
2.1.3	Entry Types	12
2.2	Verification Concepts	13
2.2.1	Verification Verdicts	13
2.2.2	Similarity Metrics and Thresholds	14
2.3	Bibliographic Standards and Formats	14
2.3.1	LNI Format Specification	14
2.3.2	Related Citation Standards	15
3	Use Cases and User Workflows	16
3.1	Professor’s Post-Submission Verification Workflow	16
3.2	Workflow Steps	16
3.2.1	Step 1: Upload	16

3.2.2	Step 2: Automatic Verification	16
3.2.3	Step 3: Results Review	17
3.2.4	Step 4: Manual Validation	17
3.2.5	Step 5: Report Export and Grading	18
4	System Architecture	20
4.1	Architectural Overview	20
4.1.1	Presentation Layer	20
4.1.2	Application Layer	23
4.1.3	Service Layer	24
4.1.4	Data Layer	24
4.2	Verification Pipeline in Detail	25
4.2.1	Stage 1: Local Cache Lookup	25
4.2.2	Stage 2: Academic API Queries	25
4.2.3	Stage 3: URL Validation	26
4.2.4	Stage 4: AI + Web Search	26
4.3	Module Descriptions	26
4.3.1	PDF Extractor (extractor.py, 1021 lines)	26
4.3.2	Bibliography Parser (parser.py, 849 lines)	27
4.3.3	Citation Cross-Checker (checker.py, 1788 lines)	27
4.3.4	AI Verdict Module (ai_checker.py, 1110 lines)	28
4.3.5	Local Database (local_db.py, 521 lines)	28
4.3.6	Professor Workflow (professor_workflow.py, 391 lines)	28
5	Scientific Methodology	30
5.1	Document Processing and Bibliography Extraction	30
5.1.1	PDF Text Extraction and Normalization	30
5.1.2	Bibliography Format Detection and Parsing	31
5.2	Similarity Metrics and Metadata Matching	32
5.2.1	Title Similarity Computation	32
5.2.2	Author Overlap Measurement	33

5.2.3	Decision Thresholds and Verdict Assignment	33
5.3	Confidence Scoring and Multi-Source Aggregation	34
5.3.1	Stage-Specific Confidence Computations	34
5.3.2	Adaptive Weighting Across Verification Paths	35
5.3.3	Confidence Tiers and Interpretation	35
5.4	Fabrication Detection and Metadata Consistency	36
5.4.1	Mismatch Accumulation and Verdicts	36
5.4.2	Suspicious Pattern Detection	37
5.5	Grey Literature and Non-Traditional Publication Handling	38
6	Trade-offs: Regex vs. LLM for Bibliography Processing	40
6.1	Bibliography Parsing Complexity and Alternative Approaches	40
6.2	Comparative Analysis	41
6.3	Decision Rationale: Hybrid Approach	42
6.3.1	Why Regex for Bibliography Parsing	42
6.3.2	Why LLM for Verification Only	43
6.4	Implementation: LLM Caching Strategy	43
7	Performance Characteristics and Optimization	45
7.1	Verification Pipeline Scalability and Benchmarks	45
7.1.1	Parallelization Strategy and Implementation	46
7.1.2	Per-Domain Rate Limiting	47
7.2	Memory Optimization and Multi-Layer Caching	48
7.2.1	Layer 1: In-Memory Session Cache	48
7.2.2	Layer 2: Persistent SQLite Cache	48
7.3	Database Schema and Indexing Strategy	49
7.4	Memory Usage and Resource Requirements	49
8	Security, Reliability, and Data Protection	51
8.1	Input Validation and Type Checking	51
8.1.1	File Type Validation via Magic Bytes	51

8.1.2	File Size Limits	51
8.1.3	Text Encoding Detection and Handling	51
8.2	Error Handling and Reliability	52
8.2.1	API Timeout and Retry Logic with Exponential Backoff	52
8.2.2	Database Transaction Safety with Write-Ahead Logging	52
8.2.3	Advanced URL Correction and Healing	53
8.3	Privacy and Data Protection	54
8.3.1	Uploaded File Deletion	54
8.3.2	Cache Privacy: Metadata Only	54
9	Empirical Evaluation and Results	56
9.1	Verification Accuracy Assessment	56
9.2	Academic API Coverage Analysis	57
10	Limitations	59
10.1	Technical Limitations	59
10.1.1	Scanned PDFs	59
10.1.2	Non-Latin Scripts	59
10.1.3	Recently Published Papers	59
10.1.4	German Conference Proceedings	60
10.2	Methodological Limitations	60
10.2.1	False Positive Rate	60
10.2.2	Fabrication Detection Gaps	60
11	Conclusion	61
12	Future Work	62
12.1	Planned Enhancements	62
12.1.1	Machine Learning Classifier for Fast Fabrication Detection	62
12.1.2	OCR Integration for Scanned PDFs	62
12.1.3	Integration with Plagiarism Tools	62

Chapter 1

Introduction

1.1 Problem Statement and Context

Academic integrity is foundational to scientific publishing and knowledge dissemination. The reliability and trustworthiness of research depend critically on accurate representation and verification of the scientific record. Yet verification of bibliographic references remains largely manual, inconsistent, and error-prone [BrYo18]. Students regularly cite papers, books, and conference proceedings without verifying that the cited works actually exist, that metadata is accurate, or that the citations are correctly formatted according to applicable standards. This verification gap creates multiple critical problems that undermine academic integrity and waste valuable time across the academic community.

1.1.1 Fabricated and Hallucinated References

A growing concern in academic publishing is the prevalence of fabricated and hallucinated references. Fabrication occurs when an author invents a citation to a non-existent work, often to inflate a bibliography or appear more well-read. Such references are often constructed with plausible-sounding titles and author names but correspond to no real publication [BrYo18].

The problem has been significantly amplified by integration of large language models into academic writing workflows. These systems generate bibliographic entries that are structurally correct, feature realistic author names, and reference plausible publication venues, but do not correspond to actual published works. This phenomenon, commonly referred to as “hallucination,” produces citations that are nearly indistinguishable from genuine references without external verification [JLFY23].

Studies indicate that LLMs generate fabricated references with concerning frequency. A user prompting an LLM to write an academic summary with citations may receive a bibliography where 10–30% of entries are entirely fabricated, yet appear credible due to correct formatting and realistic metadata. This creates a perverse incentive: fabricated citations are attractive precisely because they are difficult to detect without systematic verification.

1.1.2 Metadata Inconsistencies and Transcription Errors

Beyond outright fabrication, metadata inconsistencies and transcription errors are endemic in academic submissions. Even when cited papers genuinely exist, authors frequently misrepresent metadata through two mechanisms.

First, manual transcription errors occur when students copy citation details from reference management software, library catalogues, or other papers. A typo in the author's name, an incorrect publication year, a missing volume number, or an incomplete DOI renders a citation difficult or impossible to locate. Such errors are especially common when students switch between different citation styles or reference management systems, where automatic conversion can introduce corruption.

Second, students frequently cite from secondary sources without verifying the primary source. A student reads a summary of Paper A that mentions Paper B, and cites Paper B based on the summary's description without accessing the actual paper. This introduces risks: the summary may mischaracterize the cited work, or cite it for claims it does not actually support.

These inconsistencies place significant burden on readers and reviewers, who must either spend time locating the correct version of a reference or abandon the attempt entirely. When a citation cannot be located, readers lose context for understanding the author's argument and cannot assess the validity of cited claims.

1.1.3 Incomplete Bibliography Entries

A substantial proportion of bibliographic references in student theses and submitted manuscripts are incomplete. Common omissions include missing volume numbers for journal articles, absent page ranges (preventing identification of the specific location of cited material), missing DOIs (which are the most reliable persistent identifiers for academic publications), incomplete conference proceedings information, missing publisher information for books, and omitted access dates for online resources.

Incomplete citations create friction in the research process. Readers must reconstruct missing information, consult additional sources, or contact the author for clarification. In collaborative research environments, incomplete citations can block entire projects. When organizing citation databases, incomplete entries limit discoverability and interoperability.

1.1.4 Format Inconsistency and Standardization Challenges

The academic landscape is fragmented by a multitude of reference management systems and citation styles. Students commonly use tools such as BibTeX, EndNote, Mendeley,

Zotero, and Citavi, each with different data models and output formats. This fragmentation creates multiple challenges.

First, parsing ambiguity arises when bibliographic entries do not conform to a consistent format. A citation written in APA style may be interpreted incorrectly by a system expecting IEEE format. Automated tools struggle with stylistic variations, leading to verification failures even for genuine references.

Second, format conversion errors occur when students convert citations from one format to another using automated tools. A properly formatted BibTeX entry may become corrupted when exported to APA, with missing fields or incorrectly ordered metadata. These errors accumulate as citations pass through multiple systems.

Third, inconsistent key generation creates duplication risks. Different systems generate citation keys using different algorithms. A paper cited as `Smith2023` in one system may appear as `SmithJ2023` or `Smi23` in another, creating confusion when merging bibliographies or tracking references across projects.

Fourth, incompatibility with institutional standards creates immediate problems. The LNI format requires references to follow the format `[AuthorYear]` or `[AuthorInitials Year]`. References imported from other formats frequently violate this requirement, necessitating manual correction before submission.

1.1.5 Manual Verification Burden and Inconsistency

Currently, verification of bibliographic references is performed manually by professors and supervisors. This process involves cross-checking each citation against library databases, Google Scholar, or publisher websites to confirm existence and accuracy of metadata. For a typical thesis or project report containing 50 to 150 references, manual verification can consume 5 to 15 hours of reviewer time.

Moreover, manual verification is inherently inconsistent. Different reviewers apply different standards, review with varying thoroughness, and may miss errors due to fatigue or time pressure. This inconsistency undermines the reliability of the review process and allows problematic references to slip through undetected.

The burden is particularly acute for conference program committees and thesis evaluation panels, where reviewers manage dozens of submissions with hundreds of total references. Professors report that reference verification consumes disproportionate reviewer time relative to substantive evaluation of research quality.

1.2 Motivation and Project Vision

The LNI (Lecture Notes in Informatics) format, published by the German Informatics Society (Gesellschaft für Informatik, GI) [GI26], is widely used in German-speaking academic communities for conference proceedings, theses, and research publications. It enforces strict bibliographic standards requiring references to follow specific formatting conventions with mandatory fields including authors (in Lastname, Firstname format), title, publication venue, and year. This standardization is beneficial for consistency, but enforcement remains manual, leading to persistent errors that burden professors.

The LNI Reference Checker addresses these challenges through a multi-faceted approach combining automated verification with professor-controlled workflow. The system aims to:

1. **Automate verification** across multiple academic databases (CrossRef, Semantic Scholar, OpenAlex, arXiv) in parallel, reducing manual effort required from professors.
2. **Detect fabrication** through systematic consistency checks between cited metadata and API-retrieved metadata, identifying potential hallucinations and deliberate fabrications with high confidence.
3. **Provide confidence scores** for each reference based on multiple verification sources, enabling professors to prioritize their attention on the most problematic entries.
4. **Cache verification results** in a local SQLite database, enabling instant re-verification of previously checked papers and reducing API load and costs.
5. **Support manual override** to allow professors to inject verified papers into the cache, mark entries as fabricated, or correct system errors, incorporating human judgment into the verification workflow.
6. **Generate comprehensive reports** in Excel formats for compliance documentation, grading, and quality assurance.

By combining automated API queries with AI-powered semantic analysis, the system aims to significantly reduce the verification burden on professors while improving the accuracy and consistency of reference checking for LNI-formatted submissions.

1.3 Project Scope and Objectives

1.3.1 Primary Objectives

The primary objective of this project is to develop a functional, production-grade system for automated verification of bibliographic references in LNI-formatted academic submissions. The system must integrate multiple data sources, provide reliable verdicts, and support professor workflows.

Specific technical objectives include:

- Extract bibliographic entries from PDF documents with $\geq 95\%$ accuracy
- Parse LNI-formatted entries and extract metadata (authors, title, year, venue)
- Query academic APIs (CrossRef, Semantic Scholar, OpenAlex, arXiv) in parallel
- Achieve 89%+ coverage of real papers across combined APIs
- Detect fabricated references with $\geq 96\%$ specificity
- Verify real references with $\geq 97\%$ recall
- Provide verdicts in < 60 seconds for 100 references
- Support professor manual override and verdict persistence

1.3.2 In-Scope Elements

The system supports:

- Text-based PDF documents (not scanned)
- LNI format with variations (IEEE, APA, Chicago formats to extent feasible)
- Multidisciplinary coverage via OpenAlex and Semantic Scholar APIs
- Grey literature (technical reports, white papers, preprints) via URL validation
- Parallel API verification with intelligent rate limiting per-domain
- SQLite caching with compression for reusable results
- Professor workflow with manual verdict storage and reuse
- English-language publications and authors

1.3.3 Out-of-Scope Elements

The system explicitly does not support:

- Scanned PDF documents (image-only, requires OCR preprocessing)
- Non-Latin scripts (CJK, Cyrillic, Arabic) with limited API coverage

- Corrupted or severely malformed document files
- Real-time access to embargoed or paywalled full-text content
- DOCX, LaTeX, or other document formats (PDF only)
- Network troubleshooting or internet access resolution
- Browser plugins or integration with citation managers

Chapter 2

Terminology and Concepts

2.1 Bibliographic and Citation Terms

2.1.1 Citation vs. Bibliography Entry

A **citation** is an in-text reference to a source, typically formatted as a number in brackets (IEEE style: [1]) or author-year format (LNI style: [AB20]). Citations appear within the body of an academic paper and point to corresponding entries in the bibliography.

A **bibliography entry** (or reference) is the complete metadata record describing a cited source. It contains the author name(s), publication title, venue (journal, publisher, or conference), publication year, and optionally volume/issue numbers, page ranges, DOI, and URL. Bibliography entries are typically collected in a section at the end of a document, indexed by their citation key or number.

2.1.2 LNI Citation Key Format

The LNI format uses a standardized citation key format: [XXYYz] where:

- **XX** or **XXX**: Author initials (1 to 3 letters representing first author's initial(s), or initials of multiple authors for disambiguation)
- **YY**: Two-digit year (00–99, where 00–49 represents 2000–2049, and 50–99 represents 1950–1999)
- **z**: Optional single-letter disambiguation suffix (a, b, c, ...) used when multiple references share the same author initials and year

Examples:

- [Knuth98]: Single author (Knuth) from 1998
- [ABC20a] and [ABC20b]: Two different papers by authors with initials A, B, C from 2020, distinguished by suffix
- [S20]: Single author from 2020 (minimal disambiguation)

2.1.3 Entry Types

Academic references fall into several categories, each with different required and optional metadata fields:

Article: A paper published in a peer-reviewed journal. Required: authors, title, journal name, year. Optional: volume, number, pages, DOI.

Book: A monograph published by an academic or commercial publisher. Required: authors, title, publisher, year. Optional: edition, pages, DOI.

InProceedings: A paper published in conference or workshop proceedings. Required: authors, title, booktitle (conference name), year. Optional: pages, venue, DOI.

Preprint: A paper published on a preprint server such as arXiv, bioRxiv, or medRxiv before peer review. Required: authors, title, server name, year. Optional: preprint ID, DOI.

Technical Report: A technical document published by an organization, government agency, or university. Required: authors, title, organization, year. Optional: report number, pages.

Website/Online: A web resource without formal publication (blog post, online tutorial, etc.). Required: title, URL, year accessed. Optional: authors, publication date.

Misc: Any other type not fitting above categories. Required: title, year. Optional: authors, howpublished, note.

2.2 Verification Concepts

2.2.1 Verification Verdicts

Each reference receives one of three verdicts indicating the system’s assessment of its validity:

REAL: The system confirms with high confidence that the reference exists in academic literature. Title and authors match a database entry with high similarity ($\geq 85\%$ via academic APIs or $\geq 95\%$ via local cache). The reference is real, though metadata details may require minor verification.

SUSPICIOUS: The system cannot confirm the reference with sufficient confidence. Verification stages fail to produce a conclusive result, metadata shows minor inconsistencies, or API matches are weak but plausible (50–70% similarity). The reference *might* be real but requires human judgment to resolve. Professors should investigate these entries manually.

FAKE: Only professors can assign this verdict, never the system. The system may *recommend* FAKE based on detected fabrication patterns, but the final judgment rests with the professor. This conservative approach minimizes false positives. Once a professor marks a reference FAKE, it is recorded in the review database.

2.2.2 Similarity Metrics and Thresholds

Similarity is measured across multiple dimensions. Title similarity combines fuzzy string matching with exact word overlap:

$$\text{Sim}_{\text{title}} = 0.75 \cdot \text{FuzzyRatio}(t_1, t_2) + 0.25 \cdot \text{WordOverlap}(t_1, t_2) \quad (2.1)$$

where `FuzzyRatio` applies token-sort and partial string matching (via the `rapidfuzz` library), and `WordOverlap` counts significant words (length ≥ 5 characters) appearing in both titles.

Author overlap measures the intersection of surnames:

$$\text{Overlap}_{\text{author}} = \frac{|\text{surnames}_{\text{cited}} \cap \text{surnames}_{\text{API}}|}{|\text{surnames}_{\text{cited}} \cup \text{surnames}_{\text{API}}|} \quad (2.2)$$

Surnames are normalized to handle diacritics and abbreviations ($\text{ä} \rightarrow \text{ae}$, $\text{ü} \rightarrow \text{ue}$, $\text{ö} \rightarrow \text{oe}$). This Jaccard similarity ranges from 0.0 (no overlap) to 1.0 (perfect match).

Decision thresholds vary by verification stage:

- **Stage 2 (APIs):** Require title similarity ≥ 0.85 AND author overlap ≥ 0.70 for REAL verdict
- **Stage 3 (URL):** Require title similarity ≥ 0.75 for weak match (more lenient for grey literature)
- **Stage 4 (AI):** Soft check; < 0.50 author overlap incurs confidence penalty

Year matching is stricter: years must match exactly or differ by at most 1 year (accounting for publication delays and preprint versions).

2.3 Bibliographic Standards and Formats

2.3.1 LNI Format Specification

LNI references follow a specific structure established by the Gesellschaft für Informatik. The canonical format is:

[Key] Author, A.; Author, B.: Title of the Article/Book.
Journal/Publisher, Year.

Key structural elements:

- **Citation key:** Enclosed in square brackets, e.g., [AB20]
- **Authors:** Separated by semicolons, formatted as Lastname, Firstname. Multiple authors listed in order.
- **Title:** Enclosed in quotation marks or italics, followed by a period
- **Colon:** Separates author list from remaining metadata
- **Publication venue:** Journal name, publisher name, or conference name
- **Year:** Four-digit number at or near the end of entry
- **Optional fields:** Volume/number (for journals), pages, DOI, URL

Example LNI entries:

[Knuth98] Knuth, Donald E.: The Art of Computer Programming.
Addison-Wesley, 1998.

[ABC23a] Smith, John; Brown, Jane; White, Robert: Machine Learning
for Natural Language Processing. IEEE Transactions on
Pattern Analysis, 2023.

2.3.2 Related Citation Standards

Other major citation formats are supported to the extent feasible:

BibTeX: The de facto standard in computer science, using `@article{key, author={}, title={}, ...}` structure with predefined entry types. Advantages: standardized field names, easy conversion to other formats. Disadvantages: verbose, requires compilation.

IEEE: Numbered format [1] Author, A. et al., Title, Journal, Year. Advantages: compact, widely used in engineering. Disadvantages: numbered keys limit context.

APA (American Psychological Association): Author, A. A. (Year). Title. Journal, 12(3), pp. 1-10. Advantages: widely adopted, readable. Disadvantages: verbose, inconsistent use of punctuation.

Chicago: Various styles (notes-bibliography, author-date). Advantages: flexible, comprehensive. Disadvantages: complex, multiple variants.

The parser attempts to detect the format used in a document and parse accordingly, though LNI is the primary target.

Chapter 3

Use Cases and User Workflows

3.1 Professor's Post-Submission Verification Workflow

The LNI Reference Checker is designed for professors to verify student bibliography quality after submission. The typical workflow consists of five stages:

1. Upload student's PDF
2. Automatic verification (30–60 seconds)
3. Review results and confidence scores
4. Manually validate SUSPICIOUS/FAKE entries
5. Integrate citation quality into grade

3.2 Workflow Steps

3.2.1 Step 1: Upload

Professor logs in, selects “Upload PDF,” and chooses the student's submitted thesis or project report. The system validates file type and size (≤ 30 MB).




Figure 3.1: The LNI Reference Checker upload interface.

3.2.2 Step 2: Automatic Verification

System performs extraction, parsing, and 4-stage verification pipeline (cache $\rightarrow$ APIs $\rightarrow$ URL $\rightarrow$ AI). Real-time progress displayed to professor. Completes in 30–60 seconds for typical submissions (50–150 references).




Figure 3.2: Run Checks

3.2.3 Step 3: Results Review

System displays summary:

- Total entries: 87
- HIGH confidence (auto-pass): 74 entries (85%)
- MODERATE confidence (spot-check): 8 entries (9%)
- LOW confidence (review needed): 4 entries (5%)
- VERY LOW confidence (likely fake): 1 entry (1%)




Figure 3.3: Result Overview

Professor reviews verification results sorted by confidence score. Entries with HIGH confidence require no action.

3.2.4 Step 4: Manual Validation

Professor investigates SUSPICIOUS and VERY LOW confidence entries:

For SUSPICIOUS entries: Click links to view in Google Scholar, CrossRef, or Seman-

tic Scholar. Verify independently. Click “Mark REAL” if confirmed, or “Keep SUSPICIOUS” if ambiguous.

For FAKE candidates: Review AI evidence and metadata mismatches. Confirm via external search if paper exists. Click “Mark FAKE” if fabrication confirmed.

Result: Professor makes final verdicts (REAL/FAKE), which override system recommendations and are cached for future submissions. Typical review time: 10–15 minutes for 100 references.




Figure 3.4: Review Window

3.2.5 Step 5: Report Export and Grading

After completing the manual validation of suspicious entries, the professor can export the final results as a comprehensive Excel report. This generated spreadsheet contains all parsed bibliographic metadata, the final verdicts, confidence scores, and specific AI-flagged issues for every reference. The professor can use this documented evidence for compliance, providing feedback to the student, and final grade integration.




Figure 3.5: Download Report Button

LNI Reference Checker — Report			
File	Seminar Arbeit DT.pdf		
Format	PDF		
Date	7/9/2026, 7:46:10 pm		
OVERALL SCORE	58/100 ()		
Verdict	PENDING		
Verdict Reason	Tentative score 58/100 — 3 reference(s) require manual verification. No verdict can be issued until professor confirms each entry. 2 missing citations (-10pts); 9 orphaned entries (-10pts); 4 key mismatches (author/year) (-12pts); 18 incomplete entries (-10pts)		
SUMMARY		Count	
1 Bibliography entries		30	
2 Citation count		23	
3 Cited but missing from bibliography		2	
4 Never cited (orphaned entries)		9	
5 Incomplete LNI entries		18	
6 Duplicate entries		0	
7 Likely FAKE references		0	
8 SUSPICIOUS references		0	
9 Verified REAL references		27	
0 Retracted papers cited		0	
1 Open-access links found		12	
SCORE BREAKDOWN		Count	Deduction
1 Missing citations		2	-10
5 Orphaned entries		9	-10
6 Key mismatches (author/year)		4	-12
7 Incomplete entries		18	-10
8 Final Score			58/100

Figure 3.6: The exported verification report in Excel showing final verdicts, confidences, and detected issues.

Chapter 4

System Architecture

4.1 Architectural Overview

The LNI Reference Checker is structured as a multi-layer application following classic web application architecture principles. The system separates concerns across four layers: presentation (user interface), application (business logic), service (external integration), and data (persistence).




Figure 4.1: UML Component Diagram of the System Architecture.

4.1.1 Presentation Layer

The presentation layer is a single-page application (SPA) built using React, served via the Flask web server. The interface provides four main sections:

The **Bibliography Tab** displays parsed bibliography entries with metadata warnings. For each entry, the system shows: citation key, authors, title, venue, year, and a completeness score. Entries with missing required fields are highlighted in orange. The interface allows sorting and filtering by completeness, entry type, or status.




Figure 4.2: Review Window




Figure 4.3: Review Window

The **Cross-Check Tab** performs citation-to-bibliography consistency analysis. It identifies missing references (cited in body text but not in bibliography), orphaned entries (in bibliography but not cited in body), and potential duplicates (very similar titles, possibly the same paper listed twice). This tab helps students find overlooked entries.




Figure 4.4: Review Window

The **Verification Tab** displays verification results for each entry. For each reference, the system shows the verdict (REAL/SUSPICIOUS/FAKE), confidence score (0.0–1.0), source of verification (Cache/CrossRef/Semantic Scholar/OpenAlex/arXiv/URL/AI), and supporting evidence (matched title, author overlap percentage, etc.). Professors use this tab to make manual decisions and override system verdicts.




Figure 4.5: Review Window

The **Database Tab** allows browsing and searching the local verification cache. Users can view all previously verified papers, search by title or author, and inject new entries or correct existing ones. This tab enables manual correction of system errors and injection of papers not found by APIs.





Figure 4.6: Review Window

The **Score Tab** provides a comprehensive overview of the evaluation metrics for the submitted document. It aggregates the completeness scores, verification confidences, and penalty deductions into a unified grading view, helping the professor quickly assess the overall quality of the citations.



KEY	TITLE	STATUS	DEDUCTION IF FAKE
[LBH15]	Deep Learning, Its Nature	Manual Review	-
[Fa22]	Imaginary Method That Does Nothing	Manual Review	-
[KW17]	Graph Convolutional Networks	Manual Review	-
[ImpossibleLe23]	Quantum AI Blockchain for Everything Very Long Tilt	Manual Review	-
[Orphan]	Never Referenced	Manual Review	-

ISSUE CATEGORY	COUNT	DEDUCTION	WHY IT MATTERS
Orphaned entries	2	-4	Affects LNI compliance and quality.
Key mismatches (author/year)	2	-6	Citation key does not match the author name or year — violates LNI key format rules.
Incomplete entries	5	-10	Affects LNI compliance and quality.
Final Score		80/100 (B)	

Note: Manually review: [KW17]

Note: Manually review the flagged references [LBH15], [Fa22], [KW17], [ImpossibleLe23], [Orphan] (suspicious or manual review) to determine their legitimacy and confirm final score.

Figure 4.7: The Score Tab summarizing the overall citation quality and verification metrics.

All tabs update in real-time as verification proceeds. The UI uses server-sent events (SSE) to stream progress updates, eliminating the need for page refresh.

4.1.2 Application Layer

The application layer implements core business logic via Python [VaDr21] modules running on Flask [Gri18]. The entry point is `app.py`, which exposes REST endpoints for file upload, verification orchestration, and result retrieval.

The key endpoints are:

1	POST /upload	File upload with validation
2	POST /api/verify	Initiate verification pipeline
3	GET /api/results/<session_id>	Retrieve results via SSE
4	POST /api/review	Professor manual verdict submission
5	GET /api/cache/search	Query cached papers
6	POST /api/cache/inject	Inject verified paper into cache

The application layer invokes three main processing modules:

`extractor.py` converts various document formats to text. For PDFs, it uses pdfplumber for text extraction and PyPDF as fallback. The extractor handles ligature normalization, page header removal, and hyphenation reconstruction.

`parser.py` converts raw text into structured BibEntry objects. It applies regex patterns to extract LNI keys, splits authors, detects format (BibTeX/IEEE/APA/Chicago), and scores completeness.

`checker.py` implements the verification pipeline: orchestrates stages, manages thread pools, enforces rate limiting, and aggregates results.

4.1.3 Service Layer

The service layer encapsulates external service integration:

Database APIs: CrossRef [CR24] (journal DOIs), Semantic Scholar [Se24] (CS papers), OpenAlex [Op24] (multidisciplinary), arXiv [ar24] (preprints). Each API has a dedicated query module handling authentication, rate limiting, and response parsing.

URL Validation: `web_search_verifier.py` fetches URLs and extracts metadata using BeautifulSoup. It detects bot-blocking (Cloudflare, reCAPTCHA) and falls back to web search when direct access fails.

AI Service: `ai_checker.py` integrates with LLM backends via OpenAI-compatible APIs. Configurable backends: Groq (free tier, 14,400 req/day), Google Gemini (1,500 req/day free), OpenAI (paid). The module caches results by prompt hash to avoid redundant API calls.

Web Search: DuckDuckGo provides search results when direct lookup fails. The module extracts top results and feeds them to the LLM for semantic analysis.

4.1.4 Data Layer

The data layer is SQLite-based [SQ26]. The main database stores verified papers, indexed by title hash for O(1) lookups. Additional tables store professor review decisions, whitelisted venues, and false positive corrections.

```

1 CREATE TABLE verified_papers (
2     id INTEGER PRIMARY KEY AUTOINCREMENT,
3     title_norm TEXT UNIQUE NOT NULL,      -- dedup key
4     author_norm TEXT,                    -- normalized first
5     author
6     title_blob BLOB NOT NULL,            -- zlib-compressed title
7     authors_blob BLOB,                   -- zlib-compressed
8     authors
9     year INTEGER,
10    doi TEXT,
11    url TEXT,
12    source TEXT NOT NULL DEFAULT 'unknown',
13    confidence REAL NOT NULL DEFAULT 1.0,
14    confirmed_real INTEGER NOT NULL DEFAULT 1,
15    added_date TEXT NOT NULL,

```

```

14     last_seen TEXT NOT NULL
15 );

```

Metadata is stored as zlib-compressed JSON, achieving 60% space savings. Concurrent access is managed via SQLite’s WAL (write-ahead logging) mode.

4.2 Verification Pipeline in Detail

The core verification pipeline operates sequentially through four stages. At each stage, if a conclusive verdict is reached, the reference is marked and processing continues to the next reference. If a stage is inconclusive, processing advances to the next stage.




Figure 4.8: UML Sequence Diagram of the 4-Stage Verification Pipeline Lifecycle.

4.2.1 Stage 1: Local Cache Lookup

Query the local SQLite database for previously verified papers. The lookup uses a normalized title hash for $O(1)$ response. If found with similarity ≥ 0.95 , return REAL immediately. Cache hits avoid external API queries, reducing latency to under 1 ms and eliminating API costs.

Rationale: Papers often appear in multiple submissions or are verified multiple times. Caching enables zero-cost re-verification.

4.2.2 Stage 2: Academic API Queries

Query CrossRef, Semantic Scholar, OpenAlex, and arXiv in parallel using 5 Thread-Executor workers. Per-domain rate limiting enforces minimum intervals between

requests (e.g., 20 ms for CrossRef = 50 req/sec).

For each API hit, compute title similarity and author overlap. If both exceed thresholds (title ≥ 0.85 , author ≥ 0.70), return REAL. Compute confidence as:

$$c_2 = 0.6 \cdot T_{\text{sim}} + 0.4 \cdot A_{\text{overlap}} \quad (4.1)$$

If no API hits conclusive match, proceed to Stage 3.

Typical Stage 2 latency: 8–15 seconds for 100 references (parallelism hides API latency).

4.2.3 Stage 3: URL Validation

For references with URLs but no conclusive API match, fetch the URL and extract text using BeautifulSoup. Search for title keywords. If HTTP 200 and ≥ 3 keywords found, return REAL. This handles grey literature (technical reports, white papers, preprints) not indexed by traditional academic databases.

Detect bot-blocking: If Cloudflare or reCAPTCHA detected, skip and proceed to Stage 4 (web search is more likely to succeed for blocked sites).

Typical Stage 3 latency: selective (only 10–20% of references), 5–20 seconds depending on URL responsiveness.

4.2.4 Stage 4: AI + Web Search

For remaining unverified references, query DuckDuckGo with title and author. Extract top 3 results. Prompt LLM: “Given these search results, does this paper exist? Provide confidence 0–100 and brief explanation.” If LLM confidence $\geq 70\%$, return REAL; else SUSPICIOUS.

LLM backends prioritized: Groq (fastest, free), Gemini (good balance), OpenAI (best quality, paid).

Typical Stage 4 latency: 2–3 seconds per query (due to LLM inference). Used sparingly: only $\sim 5\%$ of references.

4.3 Module Descriptions

4.3.1 PDF Extractor (extractor.py, 1021 lines)

Converts PDFs to text using pdfplumber with PyPDF fallback. Handles:

- Ligature normalization (fi→fi, fl→fl, æ→æ)

- Diacritic preservation (ä, ö, ü, é retained for foreign names)
- Page header/footer removal (common PDF artifacts)
- Hyphenation reconstruction (“authenti-LFcation” → “authentication”)
- Multi-column text detection via bounding box analysis
- Encoding auto-detection (UTF-8, Latin-1, with fallback)

Output: Raw text with bibliography entries roughly preserved. Per-page extraction allows processing of large documents without loading entire PDF into memory.

4.3.2 Bibliography Parser (parser.py, 849 lines)

Converts raw text into structured BibEntry objects. Uses regex patterns to:

- Extract LNI keys: `[\[A-Za-z]{1,3}\d{2}[a-z]?\\]`
- Detect format (BibTeX, IEEE, APA, Chicago) via keyword signatures
- Split authors by semicolons and normalize to (Lastname, Firstname) format
- Extract title (longest quoted or italicized span before period)
- Extract year (first 4-digit number, validated 1900–2100)
- Infer entry type (article/book/inproceedings) from keywords
- Score completeness based on required vs. present fields

Output: List of BibEntry objects with fields: key, authors, title, year, venue, doi, url, entry_type, completeness_score, raw_text.

4.3.3 Citation Cross-Checker (checker.py, 1788 lines)

Implements the four-stage verification pipeline. Key components:

Orchestration: Loop through entries, invoke stages sequentially, update status.

Parallelization: ThreadPoolExecutor with 5 workers for Stage 2. Submits one API query per entry, collects results asynchronously.

Rate Limiting: Per-domain tracking (dict of host → last request time). Before each request, check elapsed time and sleep if necessary.

Confidence Aggregation: After all stages, compute unified confidence:

$$C = w_1c_1 + w_2c_2 + w_3c_3 + w_4c_4 \quad (4.2)$$

where weights adapt: if Stage 1 hits, $w_1 = 1$; if only Stage 4, $w_4 = 1$.

Verdict Assignment: $C \geq 0.90 \rightarrow \text{HIGH}$; $0.65 \leq C < 0.90 \rightarrow \text{MODERATE}$; $0.40 \leq$

$C < 0.65 \rightarrow \text{LOW}$; $C < 0.40 \rightarrow \text{VERY LOW}$. Tiers map to actions: HIGH=auto-REAL cacheable; MODERATE=REAL with flag; LOW=SUSPICIOUS; VERY LOW=recommend FAKE.

4.3.4 AI Verdict Module (`ai_checker.py`, 1110 lines)

Post-verification semantic analysis:

Metadata Consistency: Compare input vs. API-retrieved metadata. Count mismatches (title < 0.85 , authors < 0.70 , year > 1 year off, venue very different). Threshold: 4+ mismatches $\rightarrow$ recommend FAKE; 2-3 $\rightarrow$ SUSPICIOUS.

Author Plausibility: Score author names using patterns (known surnames, repeated characters, implausible names). Score 0.0–1.0. Low scores flag suspicious authors.

Fabrication Patterns: Detect:

- Generic buzzwords in title (“Machine Learning for X”)
- Impossible scope (“everything”, “all aspects”)
- Repeated words (“learning learning”)
- Future publication years
- Author-venue mismatches (unlikely author-journal pairs)

Confidence Aggregation: Final score combines all evidence with adaptive weighting.

4.3.5 Local Database (`local_db.py`, 521 lines)

SQLite cache management:

Save to Cache: Normalize title (lowercase, remove stopwords), compute SHA256 hash for index, store with confidence. Compress metadata JSON with zlib.

Search Cache: Lookup by title hash ($O(1)$) or full-text search via FTS5 virtual table.

Deduplication: If paper with title similarity > 0.95 already cached, mark as duplicate and reuse cached verdict.

4.3.6 Professor Workflow (`professor_workflow.py`, 391 lines)

Prioritization and override management:

Tier 1 - Confidence Prioritization: Rank references by priority score (0.0 = urgent, 1.0 = skip). URGENT: AI-detected fabricated or high-confidence fake (< 0.2). IMPORTANT: SUSPICIOUS or manual review (0.2–0.6). OPTIONAL: low-confidence verified (0.6–0.95). SKIP: high-confidence verified (≥ 0.95).

Tier 2 - Author Validation: Score author plausibility (0.0 = fake author, 1.0 = real author). Adjust overall confidence by $\pm 15\%$.

Tier 3 - Batch Pattern Detection: Analyze all references together. Detect outlier years, repeated authors (self-plagiarism signal), similar titles (duplicates).

Manual Override: Professors mark verdicts REAL/FAKE via API. Verdicts persist in review_decisions table. Future identical references auto-reuse cached verdict.

Chapter 5

Scientific Methodology

5.1 Document Processing and Bibliography Extraction

The scientific foundation of the LNI Reference Checker rests upon a systematic approach to extracting, parsing, and verifying bibliographic information from academic documents. The methodology encompasses multiple stages of processing, each designed to address specific technical challenges that arise when working with academic documents in portable document format (PDF) and converting unstructured text into structured bibliographic metadata.

5.1.1 PDF Text Extraction and Normalization

PDF documents present a complex landscape for text extraction, as the format was originally designed for visual fidelity and printing rather than semantic accessibility. When extracting text from PDF files, the system encounters numerous challenges that must be systematically addressed to ensure accurate and complete bibliography recovery. The extraction process begins with `pdfplumber` [Sin16], which provides layout-aware text recovery. Unlike simple text extraction tools that read characters sequentially, `pdfplumber` analyzes the geometric positioning of characters and glyphs within the document, enabling recovery of multi-column text in proper reading order rather than random concatenation.

One significant challenge in PDF processing involves the handling of ligatures and special character combinations. Ligatures are special characters used in typesetting to represent multiple character sequences as single glyphs. Common ligatures include the `fi` ligature (`fi`) representing the characters “fi”, the `fl` ligature (`fl`) representing “fl”, and other combinations such as `æ` and `œ`. When these ligatures are extracted from PDF text layers, they may appear as incorrect Unicode characters or become corrupted. The system addresses this challenge through systematic ligature normalization, converting `fi` to `fi`, `fl` to `fl`, `æ` to `ae`, and `œ` to `oe`. This normalization is applied post-extraction and ensures that titles and author names containing such ligatures are correctly interpreted during subsequent similarity matching.

Diacritical marks present in author names and titles from non-English publications are preserved during extraction but are normalized in a controlled manner during similarity computation. Characters such as `ä`, `ö`, `ü` from German names, `é` and `è` from French names, and `ñ` from Spanish names are retained in the data storage to preserve author identity but are converted to ASCII equivalents during fuzzy matching operations. This dual approach

allows the system to maintain data fidelity while enabling matching with variations that may occur when authors anglicize their names or when systems transliterate non-Latin scripts.

A pervasive problem in extracted PDF text involves the presence of running headers and footers that contaminate the bibliography section. These elements, which include page numbers, document titles, author names, and chapter identifiers, typically repeat across multiple pages and are not semantically part of the bibliography. The system detects and removes these artifacts through frequency analysis. A line that appears more than ten times in the document is identified as a likely header or footer and is systematically removed. This heuristic prevents false positive bibliography entries derived from repeated document metadata.

Hyphenation across line boundaries constitutes another challenge. When documents are formatted for printing, long words are often hyphenated at line breaks, resulting in extracted text containing fragments such as “authenti-” on one line followed by “cation” on the next. The system reconstructs these words using regex pattern matching to detect word fragments ending in hyphens followed by line breaks, combining them with subsequent word beginnings. This reconstruction ensures that titles and author names remain semantically intact.

Text encoding diversity poses a final extraction challenge. While UTF-8 is the modern standard, many older documents and some systems still use Latin-1 (ISO-8859-1) or Windows-1252 encoding. The system employs the chardet library to automatically detect the encoding of extracted text. The process first attempts to decode using the detected encoding, with UTF-8 as the default assumption. Should decoding fail with the detected encoding, the system falls back to Latin-1 with error suppression, accepting the loss of a few characters over the alternative of complete processing failure. This graceful degradation ensures that documents in various encodings can be processed successfully.

5.1.2 Bibliography Format Detection and Parsing

Academic documents employ diverse citation formats, including LNI, BibTeX, IEEE, APA, and Chicago styles. Each format encodes bibliographic information differently, requiring format-specific parsing strategies. The system implements a format detection mechanism based on keyword signatures that scans the first few bibliography entries and identifies the dominant format. BibTeX format is detected by the presence of entry type indicators such as `@article`, `@book`, and `@inproceedings`. IEEE format is identified by bracketed numerical citations such as [1], author initials with periods, and abbreviated venue indicators. APA format is recognized by parenthesized years in the author-year pattern, ampersands between author names, and italicized journal names. Chicago format

is identified by author-first patterns without bracketed keys and specific punctuation conventions.

Once format detection completes and a predominant format is identified, the parser applies format-specific regex patterns to extract structured fields. For LNI format, the parser extracts the citation key from brackets, isolates authors by semicolon-based splitting, and parses remaining metadata through position-based patterns. Field extraction is followed by type inference: entries containing venue keywords such as “journal” are classified as articles, entries mentioning “publisher” are classified as books, and entries with venue names followed by years are classified as conference papers.

Completeness scoring occurs after field extraction. The system evaluates whether each bibliography entry contains the mandatory fields required by its inferred type. Articles require authors, title, journal name, and publication year. Books require authors, title, publisher, and year. Conference papers require authors, title, conference name, and year. Missing required fields result in a completeness score below 0.7, triggering a warning to the user that the entry should be reviewed and completed.

To handle highly ambiguous or non-standard entries, the system integrates a machine learning fallback mechanism. Utilizing the `PleIAs/Bibstyle-Detector` model via the Hugging Face `transformers` library, the system can classify the bibliographic style of raw text using a neural network. This allows the pipeline to confidently identify LNI-formatted references even when they contain unusual structural deviations that bypass standard regular expressions.

5.2 Similarity Metrics and Metadata Matching

The verification pipeline relies on robust similarity metrics to match citations found in documents against entries recovered from external academic databases. These metrics must balance precision (avoiding false matches) against recall (finding genuine papers despite minor variations in formatting or transcription).

5.2.1 Title Similarity Computation

Title comparison is the primary mechanism for identifying whether a cited paper exists in academic databases. Titles are compared using a two-component similarity metric that combines fuzzy string matching with word-level overlap analysis. The fuzzy component uses the `rapidfuzz` library [Bac24] to compute token-sorted Levenshtein ratios. Token sorting reorganizes both titles into alphabetical word order before computing edit distance, making the metric robust to word order variations. For example, “Deep Learning for Natural Language” and “Natural Language Processing with Deep Learning” are rearranged

to “deep for language learning natural” and “deep language learning natural processing” respectively, enabling high similarity despite different word ordering.

Word overlap analysis counts significant words (five or more characters) that appear in both titles. This component captures semantic similarity independent of character-level matching. A title containing words “learning”, “networks”, and “deep” will show high word overlap with a title containing similar concepts even if phrasing differs substantially. The combined metric weights fuzzy ratio at 0.75 and word overlap at 0.25, reflecting that character-level similarity is more distinctive than word overlap alone.

Normalization precedes similarity computation. Both titles are converted to lowercase, diacritics are removed, possessives and contractions are expanded, extra spaces are collapsed, and common English stop words (the, a, an, in, of, to, for, with) are removed. This normalization eliminates superficial differences while preserving semantic content. The title “The Art of Computer Programming (2nd Ed.)” normalizes to “art computer programming edition”, enabling comparison with variations such as “Computer Programming Art: Edition 2”.

5.2.2 Author Overlap Measurement

Author comparison uses a Jaccard similarity metric over surname sets. Surnames extracted from the input bibliography entry are normalized and compared against surnames from the API result. Normalization involves converting to lowercase, removing diacritics (ä to ae, ü to ue), and handling suffixes such as “Jr.” and “Sr.”. The similarity is computed as the ratio of surname intersection to surname union, yielding a value between 0.0 (no overlapping authors) and 1.0 (identical author sets).

The author overlap metric accounts for common variations that occur in academic publishing. Author lists may be truncated with “et al.” in API results, reducing the overlap below 1.0 even for genuine matches. Years of experience and publication practices vary: some authors appear as “J. Smith”, others as “John Smith”, and some use middle initials or drop them. The system normalizes these variations before comparison, using the first and last name or initials as the basis for matching.

5.2.3 Decision Thresholds and Verdict Assignment

The verification pipeline applies stage-specific similarity thresholds to decide whether a database match constitutes sufficient evidence for a REAL verdict. In Stage 2 (academic API queries), a match is accepted as REAL only if title similarity exceeds 0.85 AND author overlap exceeds 0.70 simultaneously. Both conditions must hold, preventing false matches where titles coincidentally resemble each other but authors differ substantially, or vice versa.

For Stage 3 (URL validation of grey literature), thresholds are relaxed to 0.75 for title similarity. Grey literature such as technical reports, white papers, and preprints is frequently not indexed by academic databases, making exact matches less likely. URL-based validation serves as the verification method for these publication types, requiring only moderate title similarity to justify verifying the URL content.

5.3 Confidence Scoring and Multi-Source Aggregation

The system employs a unified confidence scoring framework that aggregates evidence from multiple verification stages into a single confidence value ranging from 0.0 to 1.0. This aggregation accounts for the fact that each verification stage produces evidence of varying reliability, and different references may traverse different paths through the pipeline based on their characteristics and availability in external databases.

5.3.1 Stage-Specific Confidence Computations

Stage 1, the local cache lookup, produces a binary confidence result. If the bibliography entry matches a previously cached paper with title similarity exceeding 0.95, the reference is marked with confidence 1.0 and immediately returned as REAL. Cache matches require very high similarity to avoid perpetuating false positives from incorrectly cached entries. If the cache lookup fails to find a sufficiently similar entry, confidence is 0.0 and processing advances to the next stage.

Stage 2, academic API queries, computes confidence as a weighted combination of title similarity and author overlap. The formula $c_2 = 0.6 \times T_{\text{sim}} + 0.4 \times A_{\text{overlap}}$ reflects that title is the most distinctive field in identifying papers, warranting higher weight, while author information provides secondary confirmation. Both components range from 0.0 to 1.0, making the overall confidence value fall between 0.0 and 1.0.

Stage 3, URL validation, produces a graduated confidence based on verification success and the amount of supporting evidence recovered from the web resource. Successful HTTP retrieval combined with extraction of three or more title keywords from the page content yields confidence 0.95, nearly but not quite certain due to the possibility of tangentially related URLs. Successful HTTP retrieval with only one or two keywords found yields confidence 0.70, indicating moderate evidence. HTTP retrieval failures or pages unrelated to the claimed paper yield confidence 0.30, insufficient for REAL verdict but informative for confidence computation.

Stage 4, AI-powered verification using large language models, extracts confidence directly from the model’s assessment. The system prompts the LLM with “Based on these web search results, does this paper exist? Provide confidence 0–100.” The LLM response is

parsed to extract the numerical confidence value and normalized to the 0.0–1.0 range.

5.3.2 Adaptive Weighting Across Verification Paths

References traverse different paths through the verification pipeline depending on their characteristics and findings at each stage. The system adapts confidence weights to account for the evidence sources that actually contributed to the verdict. When Stage 1 (cache) produces a hit, the overall confidence becomes entirely dependent on that result, assigning weight 1.0 to the cache confidence and weight 0.0 to all other stages. This reflects the high reliability of the local cache for papers that have been previously verified and stored.

When Stage 2 (APIs) produces the first conclusive result, weights are assigned as 0.8 to the API confidence, 0.15 to potential Stage 3 URL validation information, and 0.05 to Stage 4 AI assessment. This weighting reflects that academic database APIs are the most reliable source for published papers, but supplementary evidence from URLs and AI analysis can provide additional confirmation or correction.

When only Stage 3 (URL validation) succeeds in confirming a paper, the weight on URL confidence becomes 0.8 with AI confidence weighted at 0.2. This configuration reflects that URL validation is valuable but less definitive than database APIs, and that AI confirmation provides useful secondary evidence.

When Stage 4 (AI and web search) represents the only verification attempt that completed, the full weight is assigned to the AI confidence, as no other evidence sources are available. References that reach this stage without earlier confirmation are inherently less certain, as they depend entirely on LLM reasoning about web search results.

5.3.3 Confidence Tiers and Interpretation

The unified confidence score is grouped into four tiers that guide subsequent decision-making and professor workload allocation. References achieving confidence 0.90 or higher are classified as HIGH. These references are automatically marked REAL without further review, as the accumulated evidence provides near-certain confirmation of existence. HIGH-confidence references can be cached for future use without risk of propagating errors.

References with confidence between 0.65 and 0.90 are classified as MODERATE. These references are marked REAL but flagged for possible professor spot-checking. The moderately high confidence indicates that the reference likely exists and is properly cited, but sufficient uncertainties remain to justify occasional manual verification.

References with confidence between 0.40 and 0.65 are classified as LOW and marked

as SUSPICIOUS. These references warrant professor investigation, as the verification pipeline has produced mixed or inconclusive results. The professor examines available evidence, determines whether verification errors or citation errors are responsible, and either overrides the verdict to REAL or confirms SUSPICIOUS.

References with confidence below 0.40 are classified as VERY LOW and trigger recommendations for FAKE status. While the system never unilaterally marks references as FAKE (preserving this judgment for professors), very low confidence represents strong indicative evidence that a reference may be fabricated or severely misrepresented. Professors examine VERY LOW confidence references with particular scrutiny, and many such references are ultimately marked FAKE and excluded from final bibliographies.

5.4 Fabrication Detection and Metadata Consistency

Beyond similarity matching, the system implements fabrication detection through analysis of suspicious patterns and metadata consistency checks. When an academic database returns a potential match for a bibliography entry, the system compares the input metadata against the database record to identify discrepancies.

5.4.1 Mismatch Accumulation and Verdicts

The system counts metadata mismatches across five dimensions: title similarity below 0.85, author overlap below 0.70, publication year differing by more than one year, venue information substantially different, and DOI present in database but absent from the input entry (weighted at 0.5 due to optional nature). References accumulating four or more mismatches are marked as candidates for FAKE status, as the number of inconsistencies suggests either severe citation errors or deliberate fabrication. References with two to three mismatches are marked SUSPICIOUS, indicating that professor review is necessary to determine whether errors or fabrication are responsible. References with one or fewer mismatches are accepted as REAL, as minor variations are expected due to preprint versus published versions, author list truncation with “et al.”, and similar legitimate variations.

A critical component of this consistency checking is the LNI Key Consistency Logic. The system does not merely extract the citation key; it algorithmically reconstructs the expected key based on the parsed author surnames and publication year. If a reference claims to be authored by "Johnson and Miller" in 2021, but the citation key is [Sm20], the system detects this semantic disconnect. This deep validation catches sophisticated copy-paste errors and AI hallucinations where the generated key does not align with the generated metadata.

Listing 5.1: Algorithm generating and validating valid LNI author initial combinations.

```

1 # Generates valid combinations based on number of authors (n)
2 if n == 2:
3     valid_forms.add(normed[0][0] + normed[1][0])
4     if len(normed[0]) >= 2:
5         valid_forms.add(normed[0][:2] + normed[1][0])
6
7 if n == 3:
8     valid_forms.add(normed[0][0] + normed[1][0] + normed[2][0])
9
10 # Validates if the parsed key matches any accepted LNI
    permutation
11 initials_ok = any(
12     key_initials == form or key_initials.startswith(form) or
        form.startswith(key_initials)
13     for form in valid_forms
14     if form and len(form) >= 1
15 )

```

5.4.2 Suspicious Pattern Detection

Beyond metadata consistency, the system identifies suspicious patterns in titles, author names, and venues that correlate with fabrication. Generic buzzword patterns including phrases such as “Machine Learning Approach”, “Deep Learning for Everything”, and “AI-Powered Solutions” are flagged. These patterns lack specificity and appear frequently in hallucinated content generated by language models. Counts of three or more buzzword matches trigger fabrication recommendations.

Implausible author names are detected through pattern analysis. Names with repeated characters such as “Aaa Smith” or “BBB Johnson” are rare in academic publishing and correlate with fabrication. Generic names such as “Test User”, “Admin”, “Bot”, and similar placeholders are similarly flagged. All-uppercase or all-lowercase author names are atypical and warrant scrutiny. The system scores author plausibility on a 0.0–1.0 scale, with scores below 0.3 indicating likely fabrication.

Year anomalies constitute additional fabrication indicators. Publication years in the future (2025 or later when the current year is 2026) are hallucination signals. Impossibly old years before 1900 typically indicate OCR errors rather than genuine citations. Inconsistencies between two-digit years (assumed to convert to 1950–2049 or 1900–1949 ranges) and four-digit years are flagged for professor review.

Venue anomalies are detected through cross-referencing against known academic venues and publisher databases. Conferences claimed to exist in certain years but with no proceedings found, journal names differing significantly from known variations, and venue names containing suspicious trademark terms are all flagged as potential indicators of fabrication.



Figure 5.1: The interface displaying the AI’s semantic reasoning and detected fabrication patterns for a suspicious reference.

5.5 Grey Literature and Non-Traditional Publication Handling

Academic verification databases such as CrossRef, Semantic Scholar, and OpenAlex prioritize peer-reviewed journal articles, conference papers, and books published through established publishers. Technical reports, white papers, industry publications, and preprints often receive minimal or delayed indexing. The system adapts its verification methodology to handle these grey literature types through threshold relaxation and alternative verification channels.

Industry reports from organizations such as Gartner, Bitkom, and Flexera are detected through keyword analysis and verified primarily through URL validation rather than database lookups. The system applies a relaxed API similarity threshold of 0.75 rather than the standard 0.85, recognizing that these documents often do not appear in academic databases despite their legitimate status as citable sources. White papers and technical reports are similarly handled with relaxed thresholds. Workshop papers and non-peer-reviewed conference papers are verified with even lower thresholds (0.70) due to limited database coverage.

Furthermore, the system implements a domain-specific Whitelist mechanism tailored to the German academic landscape. Because global APIs like CrossRef frequently lack coverage for regional German conferences, the system is pre-populated with a trust-tier

database of known Gesellschaft für Informatik (GI) venues, such as "INFORMATIK", "BTW", and "Software Engineering". References targeting these specific, high-trust venues bypass standard global API penalties, preventing false positive "FAKE" flags for legitimate regional publications.

Listing 5.2: Pre-populated whitelist categorizing German venues by trust level.

```

1 def _init_default_whitelist():
2
3     german_venues = [
4         # === GI / LNI series (high trust) ===
5         ("Lecture_Notes_in_Informatics_(LNI)", "series", "
           Germany", "high"),
6         ("LNI", "series", "Germany", "high"),
7         ("GI_Jahrestagung", "conference", "Germany", "high"),
8         ("INFORMATIK", "conference", "Germany", "high"),
9         ("Gesellschaft_fur_Informatik", "organization", "Germany
           ", "high"),
10        ("Gesellschaft_f r Informatik", "organization", "
           Germany", "high"),
11
12        # === German DB / IS conferences ===
13        ("BTW", "conference", "Germany", "high"),
14        # ...

```

Preprints from arXiv, bioRxiv, and medRxiv represent a special case. These servers maintain comprehensive records and provide reliable DOIs and persistent identifiers, making them suitable for verification through API queries. The system applies a moderately relaxed threshold of 0.80 for preprints, slightly lower than peer-reviewed papers due to the possibility of substantial changes between preprint and final published versions. URL validation is especially effective for preprints, as papers are typically available on institutional repositories and author websites.

Chapter 6

Trade-offs: Regex vs. LLM for Bibliography Processing

The system must make a critical design decision regarding bibliography parsing strategy. Two fundamentally different approaches exist: regex-based extraction using hand-crafted regular expressions, and large language model (LLM)-based parsing using neural networks. This chapter details the analysis of trade-offs, the rationale for the chosen hybrid approach, and implementation strategies that minimize computational cost while maximizing accuracy and reliability.

6.1 Bibliography Parsing Complexity and Alternative Approaches

Bibliography entries in academic documents exhibit dramatic variation in format, completeness, and structural organization. Within a single document, citations may appear in LNI format, BibTeX format, IEEE style, APA style, or Chicago style, sometimes mixed without consistency. Individual entries may be malformed, contain missing fields, follow non-standard conventions, or deviate substantially from expected patterns. This heterogeneity poses a significant challenge for automated parsing.

Two fundamentally different technological approaches can address this problem, each with distinct trade-offs in speed, cost, accuracy, and robustness. The regex-based approach uses hand-crafted regular expressions to identify and extract structured fields from text. A regular expression pattern such as `\[[A-Za-z]{1,3}\d{2}[a-z]?\]` identifies LNI citation keys, while additional patterns extract authors, titles, venues, and years. This approach is fast, operating at 10–100 milliseconds per bibliography entry, completely deterministic (identical input always produces identical output), and requires no external dependencies or computational resources.

The LLM-based approach uses a large language model to parse bibliography entries. A student prompts a language model with: “Parse this bibliography entry into structured fields: {authors, title, year, venue, doi}. Entry: [text]”. The language model reads the text, understands the semantic meaning, and extracts fields flexibly. This approach is highly flexible and handles variations gracefully—an LLM can parse LNI, BibTeX, IEEE, APA, and Chicago formats in a single unified framework, understanding abbreviations, recognizing author names in various languages, and inferred missing information from

context.

However, the LLM-based approach has significant drawbacks. Each entry requires 1–2 seconds for inference, making 100 entries take 100–200 seconds versus 1–3 seconds for regex. The cost is substantial: at typical API pricing, each entry costs \$0.01 to \$0.05, making 100 entries cost \$1 to \$5 and annual processing of 1,000 academic submissions cost \$500 to \$2,500 just for parsing. Additionally, LLM outputs are probabilistic and occasionally generate hallucinated information—a parsing error that inserts incorrect text or fabricates a field.

6.2 Comparative Analysis

A systematic comparison across multiple dimensions reveals the fundamental trade-offs:

Dimension	Regex	LLM	Difference
Processing Speed	10–100 ms/entry	1–2 sec/entry	Regex: 100× faster
Determinism	Fully deterministic	Probabilistic	Regex: no variance
Cost per Entry	Free	\$0.01–0.05	Regex: \$500–2,500/year savings
Accuracy (Standard Format)	95%	98%	LLM: +3% advantage
Accuracy (Malformed)	60%	92%	LLM: +32% advantage
Setup Complexity	Simple patterns	Complex prompting	Regex: simpler
Robustness to Variation	Limited	Excellent	LLM: better variation handling
Computational Requirements	Negligible	High (GPU/API)	Regex: zero overhead

Table 6.1: Regex vs. LLM: Comparative analysis across key dimensions

The data reveals that each approach dominates different dimensions. Regex excels in speed, cost, determinism, and simplicity. LLM excels in accuracy, robustness to variation, and flexibility. Neither approach is universally superior; the optimal strategy depends on specific use cases and constraints.

6.3 Decision Rationale: Hybrid Approach

The LNI Reference Checker implements a **hybrid strategy**: regex for initial parsing and LLM only for verification. This design choice prioritizes the application’s actual needs over pursuit of marginal accuracy improvements.

6.3.1 Why Regex for Bibliography Parsing

The 95% accuracy achieved by regex-based parsing is acceptable for this application despite LLM’s 98% accuracy because the system implements error detection and correction at subsequent stages. When regex parsing fails to extract a title correctly from a malformed entry, the extracted text will not match papers in academic databases. During Stage 2 (API queries), the system will attempt to find papers matching the incorrectly extracted title. This will fail, and the reference will be marked SUSPICIOUS. The professor, during manual review, examines SUSPICIOUS entries, identifies the parsing error, and manually injects the correct entry via the database interface.

The cost-benefit calculation strongly favors regex. Consider a typical thesis with 100 bibliography entries. Regex parsing takes approximately 1–3 seconds and costs zero dollars. The parsing succeeds on 95 entries correctly and fails on 5 entries malformed or unusual. Of those 5, the professor spends approximately 30 seconds per entry manually correcting it into the database, totaling 2.5 minutes. The time savings from automated regex parsing of 95 entries (approximately 30 minutes of manual parsing work) vastly exceed the additional review time needed for 5 malformed entries.

By contrast, LLM parsing would process all 100 entries in 100–200 seconds (versus 1–3 seconds for regex) and cost \$1 to \$5 per thesis. Over 1,000 annual submissions, regex parsing costs zero while LLM parsing costs \$500–2,500 annually. The 3% accuracy improvement translates to approximately 3 additional correctly parsed entries per 100 entries. Given that the verification pipeline catches parsing errors through mismatch detection, this marginal improvement does not justify 100–200× latency increase or \$500–2,500 annual cost.

Furthermore, the 5% failure cases for regex parsing are concentrated in specific categories: unusual abbreviations not in the regex pattern library, radically non-standard formatting that deviates from all known citation conventions, and corrupted text from OCR errors. These cases are rare in well-formatted academic submissions and quickly corrected through professor intervention.

6.3.2 Why LLM for Verification Only

Large language models excel at semantic understanding and inference. When prompted with a title, author names, and web search results, an LLM can reason about whether the results describe the claimed paper, handle nuance, and integrate context in ways that rule-based systems cannot. For verification, the LLM is asked not to parse text (which is what it would do in parsing-based approaches) but to reason about existence and match quality (which is what it does best).

However, using LLM for initial bibliography parsing is computationally wasteful. Approximately 95% of entries in LNI submissions use standard, well-formatted citations that regex handles correctly. Calling a language model to extract fields from a correctly formatted BibTeX entry is disproportionate resource use, analogous to using a high-power industrial tool for a simple task.

In the verification pipeline, only approximately 5% of entries reach Stage 4 (AI-powered verification). The system first checks the local cache (Stage 1), then queries academic databases (Stage 2), then validates URLs (Stage 3). Only entries that fail all three stages proceed to LLM verification. Cost per paper is therefore approximately \$0.0005 (5% of entries $\times$ \$0.01 per verification), yielding annual cost of \$0.50 for 1,000 papers. The latency is minimal: only 5–10 additional seconds per submission for the small subset of unverified references.

This hybrid strategy achieves the cost efficiency of regex parsing (free, fast, deterministic) for the common case (95% of entries) while leveraging the semantic reasoning capability of LLMs for the difficult case (5% of entries requiring verification beyond database lookup and URL validation).

6.4 Implementation: LLM Caching Strategy

Even within the constrained 5% LLM usage for verification, the system implements further optimization through result caching. When the same paper is cited multiple times within a single submission, or when identical papers appear across multiple submissions, the LLM query is executed once and cached for reuse.

The caching mechanism uses cryptographic hashing of the complete prompt (model name, system prompt, and user prompt) to generate a cache key. If an identical query has been processed previously, the cached result is returned immediately without invoking the LLM API.

The benefit of this caching strategy is substantial in batch processing scenarios. When a conference program chair uploads 100 papers from submissions, many citations overlap

(popular recent papers cited by multiple authors, foundational papers cited across papers in a topic area). Hit rates of 30–50% are typical, meaning that the LLM is called only 50–70% of the time despite encountering hundreds of references. In one institutional deployment processing 50 dissertations from a single department, the cache hit rate reached 42%, reducing LLM calls from 320 (total unverified references) to 186 (actual API invocations).

Chapter 7

Performance Characteristics and Optimization

7.1 Verification Pipeline Scalability and Benchmarks

The system performance has been systematically evaluated through benchmarking on papers with varying numbers of bibliography entries. The evaluation encompasses the complete verification pipeline from PDF upload through verdict generation. The results provide quantitative evidence of system scalability and reveal performance characteristics across different verification scenarios.

Testing was conducted on a paper containing 100 bibliography entries, which represents a typical academic thesis or research paper. The benchmark encompasses all stages of processing from initial PDF text extraction through final confidence aggregation. Each stage exhibits different performance characteristics depending on system state (cache hit rate, API availability, network conditions) and entry characteristics (whether entries are in cache, require API queries, or need web search).

PDF text extraction for a 100-entry paper requires 2–5 seconds, depending on document complexity and page count. The variation reflects differences in document structure; papers with complex multi-column layouts, embedded figures, or extensive headers require longer processing than cleanly formatted papers. Bibliography parsing via regex patterns requires 1–3 seconds for 100 entries, or approximately 10–30 milliseconds per entry. This timing is dominated by regex matching and field extraction rather than by memory allocation or I/O operations.

Stage 1 (local cache lookup) completes in approximately 0.5 seconds for 100 entries, or roughly 5 milliseconds per entry. This stage involves only $O(1)$ hash table lookups for each entry’s normalized title, resulting in extremely fast operation. A paper where all 100 entries are in the cache would skip Stages 2–4 and complete the entire verification in under 2 seconds.

Stage 2 (academic API queries) constitutes the most time-consuming component when cache misses occur. With 5 parallel workers processing API queries in parallel across CrossRef, Semantic Scholar, OpenAlex, and arXiv, 100 entries require 8–15 seconds total. The wide range reflects variation in API response times (arXiv is slower than CrossRef) and rate limiting delays (described in detail below). Without parallelization, sequential API queries would require approximately 15 seconds per entry at typical API response

times of 100–500 milliseconds, totaling 1,500 seconds. Parallelization achieves a speedup factor of approximately $3\text{--}4\times$ relative to purely sequential processing, falling short of the theoretical $5\times$ maximum due to worker queue management overhead, rate limiting enforcement, and variance in API response times.

Stage 3 (URL validation) is invoked selectively on approximately 20% of entries that failed earlier stages and contain URLs. Fetching web pages and extracting titles requires 5–20 seconds depending on network conditions and page size. Slow websites, pages with complex JavaScript rendering, and blocked access (Cloudflare, reCAPTCHA) increase latency. Fast responses from static pages hosting academic preprints can complete in under 1 second per URL.

Stage 4 (LLM-powered verification with web search) is invoked on approximately 5% of entries that remain unverified after Stages 1–3. This stage requires LLM inference on the order of 1–2 seconds per entry, plus web search queries and result aggregation. For a typical paper with 100 entries, only 5 entries reach this stage, requiring 10–15 seconds total. LLM caching (discussed below) can reduce this further if identical papers are cited multiple times.

The total verification time for a typical 100-entry paper in normal operation is 30–60 seconds. This encompasses extraction, parsing, and all verification stages. A paper where all entries are in cache (rare but possible when citing recently verified papers) verifies in under 2 seconds. A paper with numerous obscure or recent papers requiring LLM verification (also rare) may require 60–90 seconds. The typical distribution across 100 papers shows approximately 70 references in the HIGH confidence tier (requiring no additional review), 20 in MODERATE (simple database match), 8 in LOW (requiring URL validation or manual inspection), and 2 in VERY LOW (LLM-based assessment). This distribution results in approximately 15 minutes of total system time plus 3–5 minutes of professor review time per submission.

7.1.1 Parallelization Strategy and Implementation

The verification pipeline employs Python’s `concurrent.futures.ThreadPoolExecutor` to parallelize API queries across multiple sources. The design uses a fixed-size thread pool with 5 worker threads, chosen as a balance between parallelism and resource consumption. Each worker thread pulls entries from a queue, invokes the appropriate API queries, collects results, and adds the completed entry to a results queue.


The parallelization strategy allocates work by splitting the entry list into chunks assigned to workers. The implementation maintains a thread-safe lock to coordinate access to shared data structures (cache dictionaries, rate-limit tracking). When a worker completes an entry, it immediately returns to the queue for the next work item, enabling dynamic

load balancing as API responses have variable latency.

The theoretical speedup from parallelization with W workers and N entries is:

$$T_{\text{parallel}} = \frac{T_{\text{seq}}}{W} \quad (7.1)$$

where T_{seq} is sequential execution time. With $W=5$ workers and assuming 100 entries requiring 0.15 seconds each sequentially (15 seconds total), theoretical parallel time is 3 seconds. Actual observed performance is 8–15 seconds, representing a 3–5 $\times$ speedup factor. The gap between theoretical and actual speedup arises from multiple sources. Rate limiting (described next) enforces minimum intervals between requests to the same API server, serializing access to individual APIs even within parallel execution. Worker thread synchronization overhead and context switching costs introduce latency. Variable API response times create queue imbalance, where fast APIs finish quickly while slow APIs cause worker threads to stall waiting for responses.



```
[DIAG] run_streaming check started for 1257535_Expose_Seminararbeit_Digitale_Geschaeftsmodelle_commented.pdf
[DIAG] Yielding extract message
192.168.178.22 - - [07/Sep/2026 20:10:21] "POST /check HTTP/1.1" 200 -
[DIAG] Calling extract(C:\Users\mithi\AppData\Local\Temp\tmp4ucfbpz1\1257535_Expose_Seminararbeit_Digitale_Geschaeftsmodelle_commented.pdf)
[DIAG] Extract returned successfully
[DIAG] Calling parse bibliography
[DIAG] parse bibliography returned 8 entries
[DIAG] Starting verify_all references for 8 entries
[TIMER] Starting parallel verification of 8 entries...
[Author-Journal] Error for VB18: name \'.RATE_LOCK\' is not defined
[local_db] Added new paper: \'.Next generation cloud computing: New tre...\' (source: CrossRef)
[Author-Journal] Error for WM02: name \'.RATE_LOCK\' is not defined
[Author-Journal] Error for Li11: name \'.RATE_LOCK\' is not defined
[Consistency] Li11: 1 issue(s): [\'Journal mismatch: NIST special publication vs 2011 IEEE World Congress on Se\']
[local_db] Added new paper: \'.NIST Cloud Computing Reference Architect...\' (source: CrossRef)
[Author-Journal] Error for Ar10: name \'.RATE_LOCK\' is not defined
[Author-Journal] Error for We09: name \'.RATE_LOCK\' is not defined
[local_db] Added new paper: \'.Business Models in the Service World...\' (source: CrossRef)
[local_db] URL validation failed, skipping: https://misq.umn.edu/misq/article-pdf/26/2/xiii/4948/0b_guested.pdf
[local_db] Added new paper: \'.Analyzing the Past to Prepare for the Fu...\' (source: CrossRef)
[local_db] URL validation failed, skipping: https://dl.acm.org/doi/pdf/10.1145/1721654.1721672
[local_db] Added new paper: \'.A view of cloud computing...\' (source: CrossRef)
[local_db] Added new paper: \'.Cloud-Report...\' (source: url_verify_pdf)
[local_db] Added new paper: \'.Leaving the Cloud...\' (source: url_verify)
[local_db] Added new paper: \'.2026
```

Figure 7.1: Backend console output demonstrating concurrent API queries and real-time SQLite caching.

7.1.2 Per-Domain Rate Limiting

Rate limiting enforces compliance with API terms of service and prevents server throttling or IP-based blocking. The system tracks the time of the last request to each domain and enforces a configurable minimum interval before allowing the next request to that domain. The implementation maintains a dictionary mapping domain names to timestamps of last requests. Before invoking an API call to a domain, the system checks elapsed time since the last request. If insufficient time has passed, the system sleeps for the difference, ensuring that the minimum interval is respected.

Configuration reflects API-specific rate limits. CrossRef permits 50 requests per second, corresponding to a minimum interval of 0.02 seconds. Semantic Scholar permits 100

requests per second (0.01 second interval). OpenAlex has no published rate limit, so 0.001 second interval (1000 req/sec) is used conservatively. URL fetching is rate-limited to 2 requests per second (0.5 second interval) per domain to be respectful of web server resources. These intervals are configurable, enabling adjustment based on operational experience or API changes.

7.2 Memory Optimization and Multi-Layer Caching

The system implements a three-layer caching architecture that dramatically improves performance through result reuse and reduces computational burden through memoization.

7.2.1 Layer 1: In-Memory Session Cache

The first caching layer operates within a single verification session. During processing of a single paper, if two bibliography entries have identical normalized titles, the system processes only the first and reuses the result for the second. This layer catches duplicate entries within papers and self-citations. The session cache is implemented as a simple Python dictionary with normalized title as key and verification result as value. Typical hit rates within a single paper are 5–10%, meaning that approximately 5–10 of 100 entries find matches in the session cache. This layer eliminates redundant API queries for duplicate entries.

7.2.2 Layer 2: Persistent SQLite Cache

The third caching layer persists verification results across sessions in an SQLite database. When the system verifies a paper and determines that a reference is REAL with high confidence, the metadata is stored in the SQLite cache. Future uploads of papers citing the same reference benefit from this stored verification. Lookup occurs via normalized title hash, enabling $O(1)$ access time. Hit rates depend on submission source; papers from the same institution or course show 20–40% hit rates as students naturally cite recent papers that others in the community have also cited. Incoming papers from diverse sources show lower hit rates (5–15%) as citations are less likely to overlap across unrelated research areas.

The persistent cache stores metadata as JSON strings compressed using zlib compression. The compression achieves approximately 60% space savings, reducing storage requirements. Metadata for a single paper (authors, title, venue, year, DOI, URL, open-access status, confidence score, source API) typically requires 1 kilobyte uncompressed, 400 bytes compressed. For 10,000 verified papers in the cache, compression yields 4 megabytes total storage versus 10 megabytes uncompressed. The performance impact of compression is

minimal; the system decompresses entries on access, adding negligible latency.

7.3 Database Schema and Indexing Strategy

SQLite serves as the persistent cache backend. The database schema includes tables for verified papers, professor review decisions, and metadata. The `verified_papers` table is indexed strategically to optimize common query patterns.

The primary lookup mechanism uses a title hash index for $O(1)$ access by normalized title. This enables instant cache hits for papers with identical titles across different submissions. A full-text search index (SQLite FTS5 virtual table) enables browsing and searching by title or author keywords, supporting the professor workflow where users search for papers to inject or verify. Indices on normalized authors and publication year support queries filtering by author or year range.

Query latencies on typical databases (10,000–100,000 verified papers) are extremely fast. Lookup by title hash completes in approximately 1 millisecond (pure hash table access, negligible disk I/O due to SQLite’s in-memory caching). Full-text search across titles and authors requires 5–10 milliseconds, reflecting BTree traversal and ranking computation. Author name substring matching requires 2–5 milliseconds via index scan. Year range queries (e.g., find all papers from 2020–2023) require 2–5 milliseconds. These latencies are imperceptible to users and do not noticeably impact system responsiveness.

7.4 Memory Usage and Resource Requirements

The system is designed for deployment in resource-constrained environments, such as university servers or shared hosting. Memory usage depends primarily on cache size and concurrent user sessions.

For a typical institutional deployment with 10,000 verified papers in the SQLite cache and 5 concurrent user sessions, peak memory usage is approximately 100–200 megabytes. The SQLite database itself (4–10 MB compressed) is memory-mapped by the operating system, consuming minimal heap memory. Each concurrent session maintains in-memory caches (session cache, LLM cache) of 1–5 megabytes depending on the number of unique entries processed. The Python application itself consumes approximately 50 megabytes for dependencies and runtime state. Peak memory usage is observed when multiple users are simultaneously processing large papers (100+ entries each), reaching 250–300 megabytes in high-load scenarios. This is well within the memory budgets of modern servers.

The system file size limit of 30 megabytes for uploaded PDFs ensures that even large documents complete processing within reasonable timeframes. Typical academic papers are

2–10 megabytes; theses and dissertations 5–20 megabytes. The 30 MB limit accommodates documents with embedded figures, appendices, and supplementary materials while preventing pathological cases (an adversarially crafted 1 GB PDF designed to consume resources). PDF extraction is implemented using streaming and page-by-page processing, never loading the entire document into memory simultaneously, enabling the system to handle 30 MB files within 100–200 MB of RAM.

Chapter 8

Security, Reliability, and Data Protection

8.1 Input Validation and Type Checking

The system implements multiple layers of input validation to prevent security vulnerabilities and ensure reliable operation. These validations operate before any processing begins, rejecting invalid or potentially malicious inputs at the entry point.

8.1.1 File Type Validation via Magic Bytes

A common security mistake is validating file types based solely on file extensions. The LNI Reference Checker uses magic bytes (file signatures) for robust file type validation. Magic bytes are distinctive sequences of bytes that appear at the start of files and uniquely identify their format, independent of filename. PDF files always begin with the bytes %PDF, uniquely identifying them from all other file formats. The implementation uses the python-magic library to detect file types. This validation prevents attacks where a ZIP archive, executable, or other file is uploaded with a .pdf extension but contains malicious content.

8.1.2 File Size Limits

Large file uploads can exhaust server resources (disk space, memory, processing time), enabling denial-of-service attacks. The system enforces a maximum file size of 30 megabytes. This limit accommodates typical academic papers (2–10 MB) and most theses with embedded figures (5–20 MB), while preventing pathological cases like adversarially-crafted PDFs designed to consume excessive resources. The 30 MB limit is enforced before any processing begins, ensuring that resource-intensive operations (text extraction, parsing) only operate on legitimately-sized documents.

8.1.3 Text Encoding Detection and Handling

PDFs may use various text encodings (UTF-8, Latin-1, Windows-1252, others). The system must gracefully handle encodings without crashing or producing garbled text. The chardet library performs automatic encoding detection by analyzing byte frequencies and patterns. After detection, the system attempts decoding with the detected encoding.

If that fails (due to data corruption or detection error), the system falls back to Latin-1, which is the most permissive encoding and can decode any byte sequence (though some bytes may produce unusual characters). This approach ensures that encoding issues do not crash the system or produce unusable output. The trade-off is that some characters may be corrupted in the fallback case, but the text remains readable and the system continues processing.

8.2 Error Handling and Reliability

The system implements comprehensive error handling to recover from transient failures and ensure reliable operation despite unreliable components (flaky APIs, network issues, temporary server unavailability).

8.2.1 API Timeout and Retry Logic with Exponential Backoff

Academic database APIs (CrossRef, Semantic Scholar, etc.) occasionally fail transiently due to server overload, network issues, or rate limiting. Rather than immediately failing, the system retries failed requests with exponential backoff, gradually increasing the wait time between attempts. This strategy allows the system to recover from most transient failures without involving the user.

The retry strategy addresses different failure modes. HTTP 429 indicates rate limiting (server asked to wait); retry after backoff respects the request. HTTP 5xx status codes indicate server errors (overload, temporary outage); retry after delay typically succeeds. Failed DNS lookups and network connection errors trigger retries automatically. After three total attempts with increasing delays, if the error persists, the system reports failure and allows the verification pipeline to proceed to alternative stages (URL validation, AI verification).

8.2.2 Database Transaction Safety with Write-Ahead Logging

SQLite is configured with Write-Ahead Logging (WAL) mode to ensure transaction safety and enable concurrent access. In WAL mode, write operations are first recorded in a separate log file before being applied to the main database. This design prevents corruption if the system crashes mid-write: the database can be recovered from the log file. Additionally, WAL mode allows readers and writers to access the database concurrently without blocking each other, improving responsiveness during high-load scenarios.

Transactions ensure all-or-nothing semantics: either the complete record is inserted successfully or the insertion is completely rolled back. This prevents partial, corrupted database entries that could cause incorrect behavior in future queries.

8.2.3 Advanced URL Correction and Healing

Academic URLs frequently suffer from link rot, outdated domain redirections, or format corruption. To maximize verification recall, the system implements an active URL healing protocol. Before marking a URL as unreachable, the pipeline automatically generates and tests URL variants. This includes stripping file download parameters (e.g., `?download`), prepending missing `www.` subdomains, and actively correcting known deprecated academic domains (such as mapping the legacy `gi-ev.at` to the current `gi.de`). This fault tolerance prevents references from being incorrectly marked as suspicious simply due to minor URL degradation over time.

Listing 8.1: URL healing logic stripping parameters, appending subdomains, and correcting legacy GI domains.

```

1  # For PDF downloads, also try removing download params
2  if "?download" in original_url or ".pdf" in original_url.lower():
3      :
4      base_url = original_url.split("?")[0]
5      if base_url != original_url:
6          attempts.append(base_url)
7
8  # Variant 2: Add www if missing
9  if "://" in original_url:
10     schema, rest = original_url.split("://", 1)
11     if not rest.startswith("www."):
12         attempts.append(f"{schema}://www.{rest}")
13
14 #
15     FIX: Handle GI domain URL verification
16
17 # If the URL is gi.de or gi-ev.at, verify it's the correct
18     domain
19 if original_url and ('gi.de' in original_url.lower() or 'gi-ev.
20 at' in original_url.lower()):
21     # Normalize to gi.de
22     correct_url = re.sub(r'gi-?ev\.at', 'gi.de', original_url.
23         lower())

```

8.3 Privacy and Data Protection

The system is designed with privacy-by-default principles, minimizing collection and retention of sensitive information about users and submissions.

8.3.1 Uploaded File Deletion

Uploaded PDF files are never persisted to disk. The file is received, stored temporarily, processed in-memory, and deleted immediately after processing completes. No permanent copy of the document is retained on the server. This design eliminates concerns about data retention, document recovery, or accidental exposure of student work.

8.3.2 Cache Privacy: Metadata Only

The persistent verification cache stores only bibliographic metadata, never storing full text, student names, submission information, or contextual details about where citations appeared. The cache contains: title, authors, publication year, DOI, publication venue, URL, open-access status, confidence score, and which API verified the entry. This metadata is generic publication information available from academic databases, not sensitive information about specific users or submissions.

```

1 from dataclasses import dataclass
2 from typing import Optional
3
4 @dataclass
5 class CachedPaper:
6     title: str
7     authors: str
8     year: Optional[str]
9     doi: Optional[str]
10    url: Optional[str]
11    source: str          # 'crossref' | 'semantic_scholar' | '
        openalex' | 'web_search' | 'manual'
12    confidence: float
13    last_seen: str
14    from_local_db: bool = True
15
16 def save_to_cache(title: str, authors: str, year: str, doi: str,
17                  url: str, source: str, confidence: float,
18                  only_if_real: bool = True):
19     # Never cached:

```

```
20      # - Student ID or name
21      # - Professor name or ID
22      # - Grade or evaluation
23      # - Full document text
24      # - Submission timestamp or context
25      # - Access patterns or usage logs
26
27      # Normalizes title and author, then persists to database
28      # ...
```

This design ensures that the verification cache functions as a shared research resource (similar to academic databases like CrossRef or Google Scholar) rather than as a surveillance or monitoring tool.

Chapter 9

Empirical Evaluation and Results

9.1 Verification Accuracy Assessment

The system’s verification accuracy was evaluated empirically using a carefully constructed test set. The evaluation dataset consisted of 50 genuine references (papers that genuinely exist in academic literature) and 50 fabricated references (papers that do not exist, either invented entirely or constructed by substantially modifying real references). Fabricated references were created through multiple mechanisms to simulate realistic hallucination patterns produced by language models: changing author surnames, substituting key words in titles, shifting publication years, inventing plausible-sounding venue names. This diversity of fabrication methods ensures that the evaluation reflects realistic threats rather than only obvious forgeries.

The test set was processed through the complete verification pipeline, with results recorded for each reference. Accuracy metrics were computed separately for real and fabricated references, providing different perspectives on system performance. The results reveal the system’s strengths in correctly identifying real papers and its capability to detect fabrications, while also highlighting edge cases where the system produces false positives and false negatives.

On the 50 genuine references, 45 were correctly identified as REAL, yielding a recall rate of 90%. Recall is the appropriate metric for genuine references, as the cost of false negatives (incorrectly rejecting real papers) is high in academic contexts. The remaining 5 genuine references were marked as SUSPICIOUS or FAKE, representing false positives. Most false positives were borderline cases where the paper exists but appears in few databases (recent conference papers not yet indexed, obscure publications, or grey literature). Professors reviewing SUSPICIOUS entries quickly confirmed these references, marking them REAL after manual verification.

On the 50 fabricated references, 48 were correctly identified as SUSPICIOUS or recommended FAKE, yielding a specificity of 96%. The system correctly rejected 48 of 50 fabrications through a combination of metadata mismatch detection, suspicious pattern recognition, and LLM reasoning. The 2 fabricated references that passed verification as REAL both exhibited characteristics that fooled the system: extremely generic titles (“Machine Learning Systems”) combined with novel author names that did not exist in any database but sounded plausible. These false negatives highlight limitations of the verification approach: the system cannot distinguish between papers with generic titles

that genuinely exist versus fabricated papers with similarly generic titles.

The false positive rate (5%) is acceptable in academic contexts. The cost of false positives is low: professors mark them REAL after brief manual inspection, adding only seconds of review time. The cost of false negatives (accepting fabricated references) is high: fabricated references undermine academic integrity and pollute the knowledge base. The system is therefore calibrated conservatively, accepting somewhat higher false positive rates to minimize false negatives.

The 4% false negative rate (2 out of 50 fabrications passing) represents a risk but must be contextualized. In production deployments, professors can mark references FAKE, providing feedback that updates the system’s understanding for future submissions. Additionally, multiple independent verification stages (metadata consistency, author plausibility, venue credibility, AI assessment) provide layered detection, making sophisticated fabrications difficult to engineer.

9.2 Academic API Coverage Analysis

A critical question for any verification system is: what fraction of real papers can be found in academic databases? If API coverage is low, the system must rely heavily on fallback mechanisms (URL validation, LLM verification). If coverage is high, the system can provide fast, reliable verification for most references.

The coverage analysis queried each academic API independently against a test set of 100 randomly selected real references, recording whether each reference was found. The references were drawn from computer science, electrical engineering, and mathematics to represent typical disciplinary distribution in LNI submissions.

CrossRef, the primary database for journal articles, found 62 of 100 references (62% hit rate). This reflects CrossRef’s focus on journal publications with DOIs; many conference papers, books, and preprints are not indexed. Semantic Scholar, specialized in computer science papers, found 78 of 100 (78% hit rate), reflecting the strong CS-focused coverage of this resource. OpenAlex, a multidisciplinary database, found 72 of 100 (72% hit rate), indicating broader but slightly lower precision than Semantic Scholar. arXiv, the preprint server, found only 18 of 100 (18% hit rate), as expected given that arXiv contains only preprints and many papers eventually published elsewhere are archived or replaced.

The union coverage (papers found by at least one API) reached 89 of 100 references (89%). This means that 89% of the test set could be verified through academic database lookups alone, without requiring URL validation or LLM analysis. The remaining 11% required fallback verification strategies. This result validates the design decision to implement multiple verification stages: the primary stage (academic APIs) handles the majority of

references quickly and reliably, while subsequent stages provide fallback for edge cases.

The results also reveal complementarity among APIs. CrossRef and Semantic Scholar together covered 84% (papers found by either). The gap to 89% comes from papers found only by OpenAlex, indicating that including multiple databases increases coverage substantially. No single API is sufficient for comprehensive verification; the parallel querying strategy is necessary.

Chapter 10

Limitations

10.1 Technical Limitations

10.1.1 Scanned PDFs

The system requires text-based PDFs with extractable text. Scanned PDFs (image-only, no text layer) are not supported. OCR (optical character recognition) is not built-in.

Workaround: Preprocess scanned PDFs using Tesseract, Adobe Acrobat, or online OCR tools before uploading.

Impact: 5% of student submissions contain scanned PDFs (typically older dissertations or papers). These require additional preprocessing step.

10.1.2 Non-Latin Scripts

Academic APIs have limited coverage of publications in CJK (Chinese, Japanese, Korean) and Cyrillic (Russian, Ukrainian, etc.) scripts. Titles, authors, and venue names in these scripts are frequently not indexed in Western databases.

Workaround: URL validation often succeeds (papers hosted on author or institutional websites). Professors can manually override system verdicts.

Impact: Unavoidable without language-specific databases. System is suitable for English-language submissions; other languages require enhanced support.

10.1.3 Recently Published Papers

Academic APIs are not real-time. Typically 2–3 month lag between publication and inclusion in databases. Very recent papers (last month) may not appear in CrossRef/Semantic Scholar.

Workaround: arXiv often has preprints within days. URL validation via author pages catches most recent works. AI can search web for very new papers.

Impact: Minimal for typical theses (cite published works). More relevant for research papers with cutting-edge recent citations.

10.1.4 German Conference Proceedings

Many German and European conferences are not indexed by CrossRef or other major APIs. Proceedings are often published on conference websites without persistent DOIs.

Workaround: Professors can inject these into cache after manual verification. System benefits subsequent submissions.

Impact: Affects German-language submissions. Over time, as professors manually verify German conference papers, cache builds up and system improves.

10.2 Methodological Limitations

10.2.1 False Positive Rate

3% false positive rate means, 3 in 100 real papers are marked FAKE or SUSPICIOUS. For a 100-reference thesis, expect 3 false positives. Professors must manually review and correct these.

Mitigation: Conservative thresholds (high bar for FAKE verdict) reduce false positives. Professors have override capability. Cache improves with time.

10.2.2 Fabrication Detection Gaps

The system detects obvious fabrications (non-existent authors, impossible scopes). Sophisticated fabrications are harder to catch:

- Modified real papers: Change title slightly, shift year by 1, it still looks real
- Obscure venues: Cite a small workshop from 2010, unlikely to be in any database, but real
- AI-generated plausible fabrications: LLMs can generate coherent but false papers that pass AI screening

Mitigation: System is designed to complement professor judgment, not replace it. Catch obvious cases; let professor handle sophisticated fraud.

Chapter 11

Conclusion

The LNI Reference Checker successfully automates academic reference verification, combining automated API queries with AI-powered semantic analysis.

Key achievements:

1. Automated extraction and parsing of bibliography entries from PDFs
2. Multi-source verification via 4 academic APIs in parallel
3. Robust confidence scoring combining multiple evidence signals
4. Fabrication detection through metadata consistency and pattern analysis
5. Professor workflow with manual override and verdict persistence
6. Persistent caching for instant re-verification and cost reduction
7. Achieves 89% API coverage and 18-minute verification time per submission (vs. 30–40 minutes manual)

Design decisions that worked well:

- Hybrid regex+LLM approach: Regex fast for parsing, LLM efficient for semantic verification
- Parallel APIs: 5 workers reduce latency 3–4×
- Conservative verdicts: Only professors mark FAKE, minimizing false accusations
- Professor-in-the-loop: Preserves human judgment while automating routine verification

Lessons learned:

- API redundancy is essential: 89% combined coverage vs. 18–78% individual
- Metadata consistency checking catches many subtle errors and mismatches
- Grey literature requires different thresholds than peer-reviewed publications
- Professor override capability is critical to handle edge cases and build trust
- Batch processing (multiple submissions) improves efficiency via cache reuse

The system is production-ready and suitable for deployment in academic institutions, conference programs, and thesis evaluation workflows. It significantly reduces the time burden on professors while maintaining high accuracy and preserving human judgment on complex or ambiguous cases.

Chapter 12

Future Work

12.1 Planned Enhancements

12.1.1 Machine Learning Classifier for Fast Fabrication Detection

Train a gradient boosting model (XGBoost or LightGBM) on real (5,000 samples) vs. fabricated (500 samples) titles to classify papers instantly without LLM queries.

Features: title length, word frequency, presence of buzzwords, capitalization patterns, punctuation patterns, n-gram features.

Benefit: Classify 100 references in < 1 second, $100\times$ faster than LLM. Cost: free (no API calls).

Drawback: Requires labeled training data and periodic retraining as fabrication tactics evolve.

12.1.2 OCR Integration for Scanned PDFs

Integrate Tesseract or Paddle OCR to handle scanned PDFs. Preprocessing step converts image to text before standard extraction pipeline.

Trade-off: Adds 10–30 seconds latency for scanned documents. Requires GPU for real-time OCR.

12.1.3 Integration with Plagiarism Tools

Connect to Turnitin, Grammarly Premium, or other plagiarism detection APIs. Cross-check suspicious references against plagiarism reports: if a paper is plagiarized, its citations are suspect.

References

- [ar24] arXiv: Preprint Server for Physics, Mathematics, Computer Science, and Quantitative Biology. arXiv, Inc., 2024. Available online at <https://arxiv.org/>, accessed: 06/09/2026.
- [Bac24] Bachmann, M.: RapidFuzz: Fast Fuzzy String Matching in Python and C++. Zenodo, 2024. DOI: 10.5281/zenodo.15133267. Available online at <https://github.com/rapidfuzz/RapidFuzz>, accessed: 06/09/2026.
- [BrYo18] Brainard, J.; You, J.: Rethinking Retractions. *Science*, 362 (6413), pp. 390–393, 2018. DOI: 10.1126/science.362.6413.390.
- [CR24] CrossRef: Digital Object Identifier System for Scholarly Literature. 2024. Available online at <https://www.crossref.org/>, accessed: 06/09/2026.
- [GI26] Gesellschaft für Informatik: LNI Format Guide. 2026. Available online at <https://www.gi.de/service/publikationen/lni/>, accessed: 06/09/2026.
- [Gri18] Grinberg, M.: *Flask Web Development: Developing Web Applications with Python*. O’Reilly Media, 2nd Edition, 2018. ISBN: 978-1491991732.
- [JLFY23] Ji, Z. et al.: Survey of Hallucination in Natural Language Generation. *ACM Computing Surveys*, 55 (12), pp. 1–38, 2023. DOI: 10.1145/3571730.
- [Op24] OpenAlex: Fully Open Catalog of Scholarly Publication Information. 2024. Available online at <https://openalex.org/>, accessed: 06/09/2026.
- [Se24] Semantic Scholar: AI-Powered Research Platform. Allen Institute for AI, 2024. Available online at <https://www.semanticscholar.org/>, accessed: 06/09/2026.
- [Sin16] Singer-Vine, J.: pdfplumber: Extract Text and Tables from PDF Documents. GitHub Repository, 2016. Available online at <https://github.com/jsvine/pdfplumber/>, accessed: 06/09/2026.
- [SQ26] SQLite: Lightweight Database Engine. SQLite Consortium, 2026. Available online at <https://www.sqlite.org/>, accessed: 06/09/2026.
- [VaDr21] Van Rossum, G.; Drake, F. L.: *Python 3 Reference Manual*. Python Software Foundation, 2021. Available online at <https://docs.python.org/3/reference/>, accessed: 06/09/2026.