Perfect LNI Compliance Test

This paper studies deep learning [LBH15] and neural networks [Go16]. We build on transformer architectures [VSP17]. Recent work in NLP [De18] shows improvements. Graph neural networks [KW17] are also relevant. The ResNet model [RH15a] and variants [RH15b] address deep network training. Computer vision uses [Si14]. Optimization methods like Adam [KB14] are fundamental. Regularization techniques from [SR14] prevent overfitting. Batch normalization [IS15] improves convergence. Dropout was introduced in [SHK14].

Literaturverzeichnis

[LBH15] LeCun, Yann; Bengio, Yoshua; Hinton, Geoffrey: Deep Learning. In: Nature, Vol. 521, 2015; S. 436--444.

[Go16] Goodfellow, Ian; Bengio, Yoshua; Courville, Aaron: Deep Learning. MIT Press, 2016.

[VSP17] Vaswani, Ashish; Shazeer, Noam; Parmar, Niki; Uszkoreit, Jakob; Jones, Llion; Gomez, Aidan N.; Kaiser, Lukasz; Polosukhin, Illia: Attention Is All You Need. In: Advances in Neural Information Processing Systems, 2017; S. 5998--6008.

[De18] Devlin, Jacob; Chang, Ming-Wei; Lee, Kenton; Toutanova, Kristina: BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In: Proceedings of NAACL-HLT, 2019; S. 4171--4186.

[KW17] Kipf, Thomas; Welling, Max: Semi-Supervised Classification with Graph Convolutional Networks. In: International Conference on Learning Representations, 2017; S. 1--10.

[RH15a] He, Kaiming; Zhang, Xiangyu; Ren, Shaoqing; Sun, Jian: Deep Residual Learning for Image Recognition. In: IEEE Conference on Computer Vision and Pattern Recognition, 2016; S. 770--778.

[RH15b] He, Kaiming; Zhang, Xiangyu; Ren, Shaoqing; Sun, Jian: Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification. In: IEEE International Conference on Computer Vision, 2015; S. 1026--1034.

[Si14] Simonyan, Karen; Zisserman, Andrew: Very Deep Convolutional Networks for Large-Scale Image Recognition. In: International Conference on Learning Representations, 2015; S. 1--13.

[KB14] Kingma, Diederik P.; Ba, Jimmy Lei: Adam: A Method for Stochastic Optimization. In: International Conference on Learning Representations, 2015; S. 1--15.

[SR14] Srivastava, Nitish; Hinton, Geoffrey; Krizhevsky, Alex; Sutskever, Ilya; Salakhutdinov, Ruslan: Dropout: A Simple Way to Prevent Neural Networks from Overfitting. In: Journal of Machine Learning Research, Vol. 15, 2014; S. 1929--1958.

[IS15] Ioffe, Sergey; Szegedy, Christian: Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift. In: International Conference on Machine Learning, 2015; S. 448--456.

[SHK14] Srivastava, Nitish; Hinton, Geoffrey E.; Krizhevsky, Alex; Sutskever, Ilya; Salakhutdinov, Ruslan R.: Dropout: A Simple Way to Prevent Neural Networks from Overfitting. In: The Journal of Machine Learning Research, Vol. 15, No. 1, 2014; S. 1929--1958.